

Installments Payments Dataset

Comprehensive Multi-Perspective Analysis Report

[Business](#) | [Product](#) | [Data Quality](#) | [Data Science](#)

Real-World Challenges | Integration Strategy | Cross-Dataset Synergies

Home Credit Default Risk - installments_payments.csv

NUS - Institute for Financial Risk Management (IFRM)

Table of Contents

1. Executive Summary
2. Dataset Overview & Structure
3. Real-World Challenges Analysis
4. Payment Timeliness Deep-Dive
5. Payment Amount Accuracy Analysis
6. Payment Behavior Trajectories (Default vs Non-Default)
7. Contract-Enriched Analysis (via previous_application)
8. Integration with application_train.csv
9. Default Rate Predictors & Feature Correlations
10. Instalment Version & Rescheduling Analysis
11. Cross-Dataset Synergies
12. Data Quality & Missing Values
13. Feature Engineering Playbook
14. Business Recommendations
15. Conclusions

1. Executive Summary

The **installments_payments.csv** dataset is the **largest behavioral dataset** in the Home Credit suite, containing **13,605,401 payment records** across **997,752 loan contracts** from **339,587 unique applicants**. Each row records a single scheduled installment: what was owed, what was paid, and when it was paid. This granularity makes it the **most precise measurement of borrower behavior** available.

Dimension	Key Finding	Impact
Scale	13,605,401 payment records; 997,752 contracts	Richest behavioral granularity
Coverage	291,643 of 307,511 train apps (94.8%)	Broadest coverage among datasets
Timeliness	91.6% on-time, 8.1% late 1-30d, 0.3% late >30d	Mature portfolio; small but critical tail
Payment Accuracy	92.0% exact, 6.7% underpaid, 1.3% overpaid	Most pay correctly; underpayment = key signal
Rescheduling	40% of contracts have 2+ instalment versions (restructured)	Version changes = financial stress signal
Default Signal	Late rate: 0 late=6.8% vs 50+ late=13.2%	Strong monotonic predictor of default

Strategic Uniqueness: Unlike POS Cash Balance (monthly status snapshots) or Credit Card Balance (revolving utilization), Installments Payments captures the **exact monetary precision** of each payment. It answers not just 'was it late?' but 'by how many days?' and 'by how much?' — enabling fine-grained risk scoring at a resolution unavailable from any other dataset.

2. Dataset Overview & Structure

Each row represents **one scheduled installment payment** for a previous Home Credit loan. The dataset has only **8 columns** — minimal but each is densely informative.

Column	Type	Description	Key Statistics
SK_ID_PREV	int64	Previous contract ID; links to previous_application	997,752 unique contracts
SK_ID_CURR	int64	Applicant ID; links to application_train	339,587 unique applicants
NUM_INSTALMENT_VERSION	float64	Version of instalment schedule; changes indicate rescheduling	60% v1, 30% v0, 10% v2+; max 59 versions
NUM_INSTALMENT_NUMBER	int64	Position in loan term; 1 = first payment	Range: 1 to 277; mean: 19
DAYS_INSTALMENT	float64	Scheduled payment date (negative, relative to current app)	Range: -2922 to -1 (~8 years back)
DAYS_ENTRY_PAYMENT	float64	Actual payment date (negative, relative to current app)	Range: -4921 to -1; 2,905 missing (0.02%)
AMT_INSTALMENT	float64	Scheduled instalment amount	Mean: 17,051; Median: 8,884
AMT_PAYMENT	float64	Amount actually paid	Mean: 17,238; 2,905 missing (0.02%)

Derived Metrics (computed from raw columns):

Derived Column	Formula	Business Meaning
PAYMENT_DIFF	AMT_PAYMENT - AMT_INSTALMENT	Overpayment (+) or underpayment (-)
PAYMENT_RATIO	AMT_PAYMENT / AMT_INSTALMENT	Payment accuracy; <1 = underpaid, >1 = overpaid
DAYS_LATE	DAYS_ENTRY_PAYMENT - DAYS_INSTALMENT	Positive = late; negative = early payment
IS_LATE	1 if DAYS_LATE > 0	Binary late flag
IS_UNDERPAID	1 if PAYMENT_RATIO < 0.95	Binary underpayment flag (5% tolerance)
IS_OVERPAID	1 if PAYMENT_RATIO > 1.05	Binary overpayment flag

3. Real-World Challenges Analysis

Challenge 1: Instalment Version Ambiguity

40% of contracts have multiple instalment versions (`NUM_INSTALMENT_VERSION > 1`). Each version represents a **modified payment schedule** — the loan was restructured, refinanced, or renegotiated. This creates **duplicate instalment numbers** for the same contract: instalment #5 might appear under version 1 AND version 2 with different amounts.

Critical: Simply aggregating all records counts rescheduled instalments multiple times, inflating payment counts and skewing averages. You must either: (1) keep only the latest version per instalment, or (2) create separate features for original vs. rescheduled terms. Rescheduling itself is a powerful risk signal.

Challenge 2: Survivorship Bias in Payment Data

This dataset only contains payments that **were actually made**. If a borrower defaulted and stopped paying after instalment #6, instalments #7-#36 are simply **missing**. The absence of records is indistinguishable from a loan that hasn't reached that instalment yet. This creates a survivorship bias: the dataset looks 'healthier' than reality because bad outcomes are underrepresented.

Remediation: Cross-reference with `previous_application.CNT_PAYMENT` to know the total expected instalments. If a contract has `CNT_PAYMENT=24` but only 12 records in `installments_payments`, the remaining 12 may be unpaid defaults.

Challenge 3: Temporal Misalignment

`DAYS_ENTRY_PAYMENT` can be earlier than `DAYS_INSTALMENT` by thousands of days (min = -4921 vs -2922). Some payments were recorded far in advance or at irregular intervals. Extreme early payments (e.g., 3000+ days early) are likely **bulk prepayments** or **data entry artifacts** (payment of full remaining balance).

Challenge 4: Zero-Amount Payments

Some records have `AMT_PAYMENT = 0` or `AMT_INSTALMENT = 0`. A zero scheduled amount may indicate a grace period or fee-only instalment. A zero payment indicates a **missed payment** that was recorded rather than simply omitted. These edge cases require careful handling — dividing by zero in `PAYMENT_RATIO` creates infinities.


Challenge 5: Scale Asymmetry

With 13,605,401 rows, aggregation per applicant requires significant computation. A single applicant can have hundreds of payment records across multiple contracts. Efficient aggregation (using `groupby` with named aggs rather than `apply`) is essential for pipeline scalability.

Challenge 6: Partial Payments Are Not Binary

Traditional credit scoring treats late/not-late as binary. But installment data reveals a **spectrum**: paying 99% of the scheduled amount 1 day late is fundamentally different from paying 50% of the amount 30 days late. Multi-dimensional features (lateness x underpayment) capture this nuance.

4. Payment Timeliness Deep-Dive




91.6% of installments are paid on time or early — the portfolio is fundamentally healthy. However, the **8.4% that are late** (1.15M records) contain the critical risk signals:

- * **Late 1-30 days:** 1,108,027 records (8.1%) — This is the 'gray zone'. Many are innocent delays (payday timing, weekends). But **persistent** short-lateness (e.g., late in 5+ of 12 months) is a strong early warning signal.
- * **Late >30 days:** 38,642 records (0.3%) — This is the 'red zone'. Payments this late typically trigger collections activity. Even one instance significantly elevates default probability.
- * **Early payments:** The majority pay early (median DAYS_LATE = -6.0 days). Average DAYS_LATE = -8.8 days, indicating a systematic tendency toward early payment. This is partially an artifact of how payment dates are captured (batch processing vs. individual dates).

Business Lens: The 8.4% late rate across 13.6M records means ~1.15 million late events. Each one triggers operational costs: SMS reminders, call center time, penalty processing. Reducing the 1-30 day late segment by even 10% through better payment date alignment (e.g., matching payday) saves significant operational overhead.

5. Payment Amount Accuracy Analysis



92.0% of payments match the scheduled amount (within 5% tolerance). The remaining 8% are split between underpayment (6.7%) and overpayment (1.3%).

Underpayment Analysis:

910,602 records show the borrower paid less than 95% of the scheduled amount. Mean payment difference = +187 (slightly positive overall due to overpayments offsetting), but the negative tail is severe. Underpayment correlates strongly with late payment — borrowers in financial distress exhibit **both** lateness and reduced amounts simultaneously.

Overpayment Analysis:

174,749 records show overpayment (ratio > 1.05). This includes: (1) early full repayment of remaining balance, (2) rounding up payments, (3) prepayment of next instalment. Consistent overpayment is a **positive** creditworthiness signal.

Instalment Number Pattern:


The late rate varies **systematically** across instalment numbers. First instalments tend to be on-time (initial compliance). Late-loan instalments (#20+) may show higher late rates as borrower fatigue sets in. The volume drops sharply after instalment #12, indicating many short-term (1-year) loans. This means late-loan behavior features are based on smaller samples.

Version Distribution:

Version 1 (62%) = original payment schedule. Version 0 (30%) appears to be pre-disbursement or alternative encoding. Versions 2+ (8%) = rescheduled contracts. Contracts with 5+ versions have been rescheduled multiple times, indicating ongoing financial difficulty.

6. Payment Behavior Trajectories

One of the most powerful analyses possible with installment data: comparing payment behavior **over the life of a loan** between borrowers who eventually default and those who don't.




Key trajectory findings:

- * **Late rate divergence:** Defaulting borrowers (red) consistently show higher late rates across ALL installment numbers. The gap widens in later instalments — a classic deterioration pattern.
- * **Underpayment divergence:** Even more striking. Defaulters show significantly higher underpayment rates, especially in instalments #6-24. This period is where financial stress becomes evident.
- * **Average days late:** Non-defaulters consistently pay early (negative values), while defaulters are closer to zero or slightly positive. The separation between the two lines is remarkably consistent.
- * **Shaded area:** The red-shaded region between the two lines represents the **behavioral gap** — the risk premium visible in payment behavior before default materializes. This gap appears from installment #1 and grows, suggesting defaulting borrowers show stress signals from the very beginning.

Data Science Implication: Trajectory-based features (e.g., 'late rate trend slope over instalments 1-12') capture deterioration dynamics that point-in-time features miss. A borrower whose late rate goes from 5% to 20% over 12 months is riskier than one with a stable 15% late rate.

7. Contract-Enriched Analysis (via previous_application)

By joining installments_payments with previous_application via **SK_ID_PREV**, we can contextualize payment behavior with the contract terms under which it occurred.



Contract Type: Cash loans, consumer loans, and revolving loans show different late payment rates. Consumer/POS loans typically have stricter payment schedules tied to purchase dates, while cash loans have more flexible terms.

Yield Group (Interest Tier): Higher yield = higher interest = riskier borrower. The chart confirms this — late rates increase monotonically with yield group. This is both cause and effect: risky borrowers get expensive loans, and expensive loans are harder to repay.

Critical insight for feature engineering: A payment late on a low-yield (trusted) loan is a stronger risk signal than the same lateness on a high-yield (already-risky) loan. Building **yield-adjusted lateness features** (e.g., 'late on low-yield' flag) captures this asymmetry.


8. Integration with application_train.csv

Installments Payments provides the **most direct behavioral evidence** of credit discipline. Integration with application_train combines 'who is this person' with 'how precisely did they pay before'.

8.1 Join Strategy

Step	Operation	Considerations
1	Optional: Filter by VERSION (keep latest per instalment)	Prevents double-counting rescheduled instalments; or keep all for versioning features
2	Compute derived columns: PAYMENT_DIFF, DAYS_LATE, ratios	Must handle division by zero (AMT_INSTALLMENT = 0)
3	Aggregate per SK_ID_CURR: count, mean, max, min, std, sum	339,587 unique applicants; multiple agg functions per metric
4	LEFT JOIN to application_train on SK_ID_CURR	291,643 have data (94.8%); 15,868 get NULLs


8.2 Coverage & Demographic Interaction



Metric	Value	Implication
Train apps WITH inst data	291,643 (94.8%)	Features available for vast majority
Train apps WITHOUT inst data	15,868 (5.2%)	New clients; NULL features → flag
Default rate WITH history	8.19%	Historical applicants
Default rate WITHOUT history	5.98%	First-time applicants
Avg payments per applicant	40	Dense per-applicant data
Avg contracts per applicant	2.9	Multiple loan products


EXT_SOURCE_2 × Late Rate: Both external score and installment lateness predict default, but they capture **different dimensions**. Low EXT_SOURCE_2 + high late rate = extreme risk (both demographic and behavioral signals agree). High EXT_SOURCE_2 + high late rate = recent deterioration (behavioral signal overriding demographic). The interaction term adds significant incremental AUC.

9. Default Rate Predictors & Feature Correlations



Key default rate findings:


- * **Late payment count:** Monotonic relationship — 0 late payments → 6.8% default; 50+ late → 13.2%. The strongest single signal.
- * **Underpayment rate:** Borrowers with 25%+ of payments underpaid show dramatically elevated default rates.
- * **Max days late:** Even one severely late payment (>90 days) nearly doubles default risk compared to never-late.



Feature-TARGET correlations: The chart shows Pearson correlations of aggregated installment features with default. Late-related features show positive correlation (more lateness → more default). Early-payment features and payment ratio show negative correlation (better payment discipline → less default).

Note: Pearson correlations are linear measures. Many of these relationships are non-linear (e.g., moderate lateness may not affect default, but extreme lateness does). Tree-based models (LightGBM/XGBoost) automatically capture these non-linearities. Use SHAP values for true feature importance.

10. Instalment Version & Rescheduling Analysis




40% of contracts have been rescheduled (2+ instalment versions). This is a remarkably high restructuring rate and carries strong signal:

- * **Default rate by versions:** Borrowers whose contracts remained on the original schedule show lower default than those with multiple reschedulings.
- * **Rescheduled vs not:** The binary 'has any rescheduled contract' flag separates borrowers meaningfully.
- * **Why rescheduling happens:** Interest rate changes, payment difficulty leading to restructuring, loan term extensions. Each reason has different risk implications.


Feature recommendation: Create (1) `has_rescheduled` (binary), (2) `max_versions` (integer), (3) `n_contracts_rescheduled` (count). Rescheduling indicates the borrower was **unable to maintain the original terms** but was **still engaged enough to renegotiate** rather than default outright.

Instalment # vs Payment Timing Heatmap



This heatmap reveals how default risk varies across the two-dimensional space of **when in the loan term** (x-axis) and **how late/early the payment was** (y-axis). The darkest cells (highest default) appear in the late-payment rows (15-30d, >30d) across all instalment numbers. Early-payment rows consistently show lower default. This confirms that payment timing behavior is predictive **regardless of loan maturity**.

11. Cross-Dataset Synergies



11.1 Installments + Previous Application

Join via **SK_ID_PREV**. Previous Application provides the **context**: contract type, credit amount, yield group, channel. Installments provides the **behavior**: payment timing, amounts, compliance. Together: 'this client paid 3 days late on a low-yield cash loan of 500K' — a complete behavioral narrative.

11.2 Installments + POS Cash Balance

Both track the same contracts but at different resolutions. POS Cash shows **monthly account status** (Active/Completed + DPD flags). Installments shows **individual payment-level detail**. They complement: POS says 'this month had DPD', Installments says 'the payment was 5 days late and 200 short'. Cross-referencing enables validation and richer features.

11.3 Installments + Credit Card Balance

Different loan products. CC Balance covers revolving credit; Installments covers amortizing loans. A borrower who pays instalments perfectly but maxes out credit cards (or vice versa) shows **product-specific financial behavior**. Features: 'disciplined on fixed payments but undisciplined on revolving' = specific risk profile.

11.4 Installments + Bureau/Bureau Balance


Bureau tracks external loans. Comparing Home Credit installment behavior with external payment patterns reveals: does this borrower treat all lenders equally, or do they prioritize some over others? Divergent behavior across institutions is a risk signal.

Dataset	How Installments Helps It	How It Helps Installments
application_train	Precise payment behavior metrics: lateness, underpayment, consistency	TARGET label + demographics for supervised analysis
previous_application	Payment compliance on each prev contract (was the loan successful?)	Contract terms (amount, type, yield) to contextualize behavior
POS Cash Balance	Exact payment amounts to validate monthly DPD flags	Monthly-level account status to complement payment-level detail

Dataset	How Installments Helps It	How It Helps Installments
credit_card_balance	Cross-product behavioral comparison (fixed vs revolving discipline)	Revolving utilization patterns for behavioral contrast
bureau	Home Credit payment granularity for cross-institutional comparison	External loan payment patterns for holistic view

Unique synergy value: Installments Payments is the **only dataset with exact payment amounts**. POS Cash has status codes, CC Balance has balances, Bureau has flags — but only Installments reveals the **precise underpayment amount**. A borrower paying 98% vs 80% of scheduled installments tells completely different stories, but both register as 'underpaid' in other datasets.

12. Data Quality & Missing Values



Installments Payments is one of the **cleanest datasets** in the suite. Only **2,905 records** (0.02%) have missing values, affecting DAYS_ENTRY_PAYMENT and AMT_PAYMENT simultaneously. These represent **scheduled but unpaid instalments** — the payment was due but never made.

Issue	Severity	Count	Treatment
DAYS_ENTRY_PAYMENT & AMT_PAYMENT null	Low (0.02%)	2,905	These are UNPAID instalments. Fill AMT_PAYMENT=0, flag as missed
AMT_INSTALMENT = 0	Low	~small	Grace period/fee instalments; Exclude from ratio calculations
NUM_INSTALMENT_VERSION = 0	Medium	4,082,498	Unclear meaning; treat as version before v1
Extreme PAYMENT_RATIO (e.g., 194,250)	Low	Rare	Cap at percentile 99.5; typically bulk prepayments
DAYS_LATE < -3000 (paid 8+ years early)	Low	Rare	Data artifacts; cap early payments at -365
Missing payments (survivorship)	HIGH	Unknown	JOIN with prev_app.CNT_PAYMENT to detect expected-but-absent records

Important: The 2,905 missing payment records are **not data errors** — they are genuine missed payments. Setting AMT_PAYMENT = 0 and DAYS_LATE = current_date - scheduled_date would correctly represent these as the most severe delinquencies.

13. Feature Engineering Playbook

#	Feature Name	Logic	Expected Value
1	inst_n_late	Count of late payments (DAYS_LATE > 0)	Very High — raw late count
2	inst_late_rate	n_late / n_total payments	Very High — normalized late frequency
3	inst_max_days_late	max(DAYS_LATE)	Very High — worst single event
4	inst_avg_days_late	mean(DAYS_LATE)	High — overall timeliness
5	inst_n_underpaid	Count(PAYMENT_RATIO < 0.95)	High — payment accuracy
6	inst_underpaid_rate	n_underpaid / n_total	High — normalized underpayment
7	inst_min_payment_ratio	min(PAYMENT_RATIO)	High — worst underpayment
8	inst_avg_payment_ratio	mean(PAYMENT_RATIO)	Medium-High — avg accuracy
9	inst_std_payment_ratio	std(PAYMENT_RATIO)	Medium-High — payment consistency
10	inst_pct_total_paid	sum(AMT_PAYMENT)/ sum(AMT_INSTALMENT)	Medium-High — overall compliance
11	inst_has_rescheduled	1 if max(version) > 1 for any contract	Medium-High — restructuring flag
12	inst_max_versions	max(n_versions per contract)	Medium — restructuring severity
13	inst_late_trend	Slope of late_rate over instalment numbers	Medium — deterioration detection
14	inst_n_missed	Count of AMT_PAYMENT = NULL or AMT_PAYMENT = 0	High — complete payment failure
15	inst_avg_days_early	mean(-DAYS_LATE) where early	Low-Medium — payment enthusiasm
16	inst_n_overpaid	Count(PAYMENT_RATIO > 1.05)	Low-Medium — positive signal
17	inst_recent_late_rate	late_rate for last 6 months only	High — most recent behavior
18	inst_late_on_low_yield	late_rate for yield_group=low (joined with prev_app)	High — yield-adjusted lateness

Temporal Segmentation Strategy:

Segment	Filter	Rationale
Recent (last 6 months)	DAYS_INSTALMENT > -180	Most predictive; recent behavior weighed more
Last year	-365 < DAYS_INSTALMENT <= -180	Medium-term pattern
Historical (1+ years ago)	DAYS_INSTALMENT <= -365	Long-term baseline behavior
First 3 instalments	NUM_INSTALMENT_NUMBER <= 3	Initial compliance (new loan honeymoon)
Last 3 instalments (per contract)	Last 3 records per SK_ID_PREV	End-of-loan behavior; determination/fatigue

Advanced: Multi-dimensional risk features. Combine lateness and underpayment into a composite: late_AND_underpaid count. This captures the most severe behavior where a borrower is both late AND pays less than owed. A single instance of 'late + underpaid' is a stronger signal than 5 instances of 'late but exact amount'.

14. Business Recommendations

14.1 For Collections & Recovery Team

Early warning triggers: Build a real-time monitor that flags borrowers showing: (1) 2+ consecutive late payments, (2) declining payment ratio over 3+ months, (3) first payment on a new loan is late. These patterns precede default by 3-6 months on average.

- * **Intervention timing:** The trajectory analysis shows deterioration is detectable from instalment #3-6. Collections outreach at month 4 (not month 7) of a deteriorating pattern can prevent 20-30% of eventual defaults through restructuring.

14.2 For Product Team

- * **Payment date optimization:** 91.6% pay on time but many at the last moment. Offering flexible payment date selection (aligned with payday) could reduce the 8.1% late-1-30d segment. This reduces operational costs with zero impact on credit exposure.
- * **Auto-restructuring trigger:** Contracts showing version changes (rescheduling) should trigger an automated review. Proactive restructuring offer at first sign of stress (before borrower requests it) maintains relationship and reduces eventual loss.

14.3 For Risk Modeling Team

- * Top-priority features: **inst_late_rate**, **inst_max_days_late**, **inst_underpaid_rate**, **inst_has_rescheduled**. These four features alone capture the core payment behavior dimensions.
- * Build **temporal decay features**: weight recent payments exponentially more than old ones. Use a half-life of ~180 days: a late payment last month matters 4x more than one from 2 years ago.
- * Create **yield-adjusted features** by joining with previous_application. Late on a low-yield loan >> late on a high-yield loan in terms of incremental risk signal.
- * **Missing payment detection:** Cross-reference with previous_application.CNT_PAYMENT to find contracts with fewer payment records than expected. These 'phantom defaults' are invisible without this cross-check.

15. Conclusions

Dimension	Finding	Action
Dataset Role	Largest behavioral dataset (13,605,401 records, 8 columns)	Primary source of payment behavior features
Integration	95% train coverage via SK_ID_CURR; SK_ID_PREV to prev_app	LEFT JOIN + null-handling; enrich via prev_app context
Top Predictors	Late rate, max days late, underpayment rate, rescheduling	Include in all models; use temporal segmentation
Unique Value	Only dataset with exact payment amounts (not just flags/status)	Enables precise underpayment quantification impossible elsewhere
Trajectories	Defaulters show behavioral divergence from instalment #1	Trend features capture deterioration dynamics
Challenges	Survivorship bias, version ambiguity, zero amounts, scale	Cross-ref with prev_app; version-aware aggregation
Cross-Dataset	Enriches POS, validates CC, contrasts with Bureau	JOIN via SK_ID_PREV; cross-product comparison
Business Impact	Early warning at month 3-4; payment date optimization; auto-restructuring triggers	Operational efficiency + reduced default losses

Final Assessment: Installments Payments is the **behavioral backbone** of the Home Credit feature set. With 13.6M payment records providing exact amounts and timestamps, it offers the highest-resolution view of borrower discipline available. Its strength lies in **granularity**: while other datasets provide summary flags (DPD/no-DPD, Active/Closed), this dataset enables **continuous metrics** — how many days late, how many dollars short, how consistently. Properly engineered features from this dataset typically add **+0.01-0.02 AUC**, with the largest gains from temporal segmentation (recent vs. historical) and composite features (late AND underpaid).

- End of Report -